

# CS695 Assignment2

Arohan Hazarika

February 2025

## 1 Flowchart (Part 1)

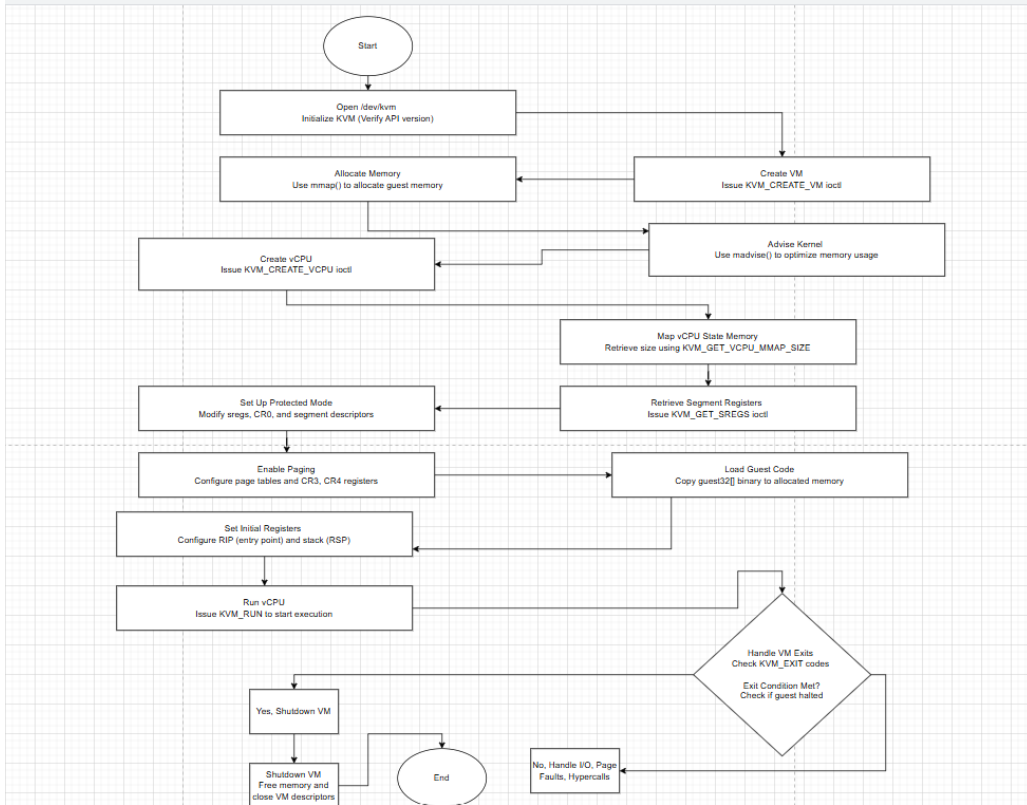


Figure 1: Flowchart explaining execution of VM in protected mode

## 2 Explanation of Code Snippets

### 2.1 (a)

```
extern const unsigned char guest64[], guest64_end[];
```

These external variables mark the start and end of the guest code binary, stored in memory. The `guest64` binary is loaded into VM memory (`vm->mem`) to allow execution inside the guest environment.

### 2.2 (b)

```
pml4[0] = PDE64_PRESENT | PDE64_RW | PDE64_USER | pdpt_addr;
```

```
pdpt[0] = PDE64_PRESENT | PDE64_RW | PDE64_USER | pd_addr;
```

```
pd[0] = PDE64_PRESENT | PDE64_RW | PDE64_USER | PDE64_PS;
```

```
sregs->cr3 = pml4_addr;
```

```
sregs->cr4 = CR4_PAE;
```

```
sregs->cr0 = CRO_PE | CRO_MP | CRO_ET | CRO_NE | CRO_WP | CRO_AM | CRO_PG;
```

```
sregs->efer = EFER_LME | EFER_LMA;
```

This code sets up a 4-level page table for memory management in long mode, configuring the **\*\*PML4\*\***, **\*\*PDPT\*\***, and **\*\*Page Directory\*\*** to establish a virtual address space. The **CR3** register is assigned the address of the top-level page table (PML4), allowing address translation. **CR0** and **CR4** enable paging and protected mode, while **EFER** enables long mode execution.

## 2.3 (c)

```
vm->mem = mmap(NULL, mem_size, PROT_READ | PROT_WRITE,  
               MAP_PRIVATE | MAP_ANONYMOUS | MAP_NORESERVE, -1, 0);  
madvise(vm->mem, mem_size, MADV_MERGEABLE);
```

The `mmap()` function is used to allocate memory for the virtual machine, mapping an anonymous memory region that serves as the VM's physical memory. The `MADV_MERGEABLE` flag in `madvise()` allows the kernel to optimize memory by merging identical pages, thus reducing redundancy and improving efficiency.

## 2.4 (d)

```
case KVM_EXIT_IO:  
if (vcpu->kvm_run->io.direction == KVM_EXIT_IO_OUT &&  
    vcpu->kvm_run->io.port == 0xE9) {  
    char *p = (char *)vcpu->kvm_run;  
    fwrite(p + vcpu->kvm_run->io.data_offset,  
          vcpu->kvm_run->io.size, 1, stdout);  
    fflush(stdout);  
    continue;  
}
```

When a guest tries to write to the I/O port `0xE9`, it triggers a VM exit (`KVM_EXIT_IO`). The hypervisor detects the exit, retrieves the data from the guest memory, and prints it to `stdout`. This mechanism allows communication between the guest and the host, simulating hypercalls or debugging output.

## 2.5 (e)

```
memcpy(&memval, &vm->mem[0x400], sz);
```

This snippet accesses memory inside the guest at address `0x400` and copies `sz` bytes into `memval`.